

Personalized Diet Recommendation System

Team – OkayNoted

Ansh Avi Khanna (IMT2021038)

Arjun Subhedar (IMT2021069)

Ritik Kumar Gupta (IMT2021098)

Pranav Bhutada (IMT2021105)



[Github](#)



Our Motivation

1.

Making Healthy Eating Easy

We're here to make it simple for everyone to eat well by suggesting personalized meals.

3.

Meeting Different Goals:

Whether it's building muscle or keeping your heart healthy, we've got you covered with tailored recommendations.

2.

Helping You Take Charge:

Our project gives you the tools to manage your health and make informed food choices.

4.

Building a Healthier Future:


By promoting balanced eating and sustainable habits, we're working towards a healthier tomorrow for everyone.




Introduction



- Eating healthily is a challenge in today's world due to busy lifestyles and conflicting dietary advice.
- Our project aims to simplify this by **offering personalized diet recommendations** for breakfast, lunch, and dinner.
- By considering specific user attributes like weight, height, gender as well as individual goals like weight loss, muscle gain, or heart health, we tailor meal plans to meet specific needs.
- With our system, users can make informed choices about their diet, leading to better health outcomes.



Dataset Creation & Scraping



Food.com - Recipes and Reviews

Data on over 500,000 recipes and 1,400,000 reviews from Food.com

[Data Card](#) [Code \(18\)](#) [Discussion \(2\)](#) [Suggestions \(0\)](#)

About Dataset

Context

The recipes dataset contains 522,517 recipes from 312 different categories. This dataset provides information about each recipe like cooking times, servings, ingredients, nutrition, instructions, and more.

The reviews dataset contains 1,401,982 reviews from 271,907 different users. This dataset provides information about the author, rating, review text, and more.

[Food.com](https://www.food.com)



Initial Dataset Exploration:

- Initially, we attempted to utilize datasets from both Kaggle and Food.com.
- Encountered challenges with a significant number of NaN values in the dataset.
- Identified difficulties in categorizing meals into distinct categories like breakfast, lunch, and dinner due to the lack of clear labeling.



allrecipes

Find a recipe or ingredient

Log In

Our Weekend Picks →



CHICKEN
Enchiladas Verdes
★★★★☆ 231 Ratings



TOMATO SALSA RECIPES
Pico De Gallo de Alicia
★★★★☆ 36 Ratings




MEXICAN
Beer Margaritas
★★★★★ 1,215 Ratings

[AllRecipes.com](https://www.allrecipes.com)

Alternative Approach: Scraping from AllRecipes.com:

- As an alternative, we decided to scrape data directly from AllRecipes.com
- This approach provided us with a superior dataset characterized by clearer meal labels and a reduction in missing data instances.



Dataset Scraping

- We utilized the Python library **Beautiful Soup** for scraping data from **AllRecipes.com**.
- Breakfast comprises 1,132 links, lunch comprises 480 links, and dinner comprises 840 links on AllRecipes.com.
- The scraping process took approximately 1.5 – 2 hrs.
- Finally, we created CSV files named **breakfast.csv**, **lunch.csv**, and **dinner.csv**.



```
def get_recipe_links(url):  
    response = requests.get(url, headers={"User-Agent": "Mozilla/  
    html_content = response.text  
    soup = BeautifulSoup(html_content, 'html.parser')  
    elements_with_class = soup.find_all(class_='comp tax-sc_rec  
    recipe_links = []  
    for element in elements_with_class:  
        links = element.find_all('a')  
        for link in links:  
            href = link.get('href')  
            recipe_links.append(href)  
    return recipe_links
```

```
https://www.allrecipes.com/recipe/260452/greek-inspired-chicken-salad/  
https://www.allrecipes.com/recipe/13637/three-bean-salad/  
https://www.allrecipes.com/recipe/13933/black-bean-and-corn-salad-ii/  
https://www.allrecipes.com/recipe/253157/weeknight-sillet-slaw/  
https://www.allrecipes.com/recipe/90500/beet-salad-with-goat-cheese/  
https://www.allrecipes.com/recipe/228283/ham-salad-spread/  
https://www.allrecipes.com/recipe/234949/authentic-russian-salad-olivye/  
https://www.allrecipes.com/recipe/14452/green-salad/  
https://www.allrecipes.com/recipe/228126/caprese-salad-with-balsamic-reduction/  
https://www.allrecipes.com/recipe/222750/quick-italian-pasta-salad/  
https://www.allrecipes.com/recipe/14439/frog-eye-salad/  
https://www.allrecipes.com/recipe/14276/strawberry-spinach-salad-i/  
https://www.allrecipes.com/recipe/222728/refreshing-cucumber-watermelon-salad/  
ritik@ritiik:~/Desktop/recsys$  
* History restored
```

```
,RecipeID,Name,Ingredients,Servings Per Recipe,Calories,  
0,1,Delicious Raspberry Oatmeal Cookie Bars,"1 cup all-p  
1,2,Quinoa and Black Beans,"1 teaspoon vegetable oil 1  
2,3,Apple Chips,"2 Golden Delicious apples, cored and t  
3,4,Peanut Butter Noodles,"8 ounces Udon noodles ½ cup o  
4,5,"Chicken, Rice, and Vegetable Soup","5 cups water, o  
5,6,Granola Bars, cooking spray 2 cups rolled oats ½ cu  
6,7,Turkey Wraps,"1 (8 ounce) package cream cheese with  
7,8,Incredibly Easy Vegetarian Chili,"1 tablespoon veget
```


Dataset Scraping

These are the nutrition data available on the site, which best fits our needs. We scraped this data and included it as columns in our dataframe.

Philosophy: Our main idea was to keep things real. We wanted to use data that reflects how people actually eat every day. By getting information from AllRecipes.com, we aimed to give practical advice that fits real-life eating habits. This way, our recommendations could be more helpful and relevant to users.



Nutrition Facts (per serving)

| | | | |
|-----------|-----------|-----------|-----------|
| 82 | 7g | 5g | 1g |
| Calories | Fat | Carbs | Protein |

[Hide Full Nutrition Label](#)

Nutrition Facts


Servings Per Recipe: 20

Calories: 82

| | % Daily Value * |
|-------------------------------|-----------------|
| Total Fat: 7g | 9% |
| Saturated Fat: 1g | 6% |
| Sodium: 5mg | 0% |
| Total Carbohydrate: 5g | 2% |
| Dietary Fiber: 4g | 13% |
| Total Sugars: 1g | |
| Protein: 1g | 2% |
| Vitamin C: 7mg | 7% |
| Calcium: 10mg | 1% |
| Iron: 0mg | 2% |
| Potassium: 262mg | 6% |

Approach - 1

Contextual Multi-Armed Bandits, Linear
UCB, Epsilon-Greedy




Data Preprocessing



Dataset Preprocessing

- **Datasets:** We have three datasets for breakfast, lunch and dinner recipes. The columns in these datasets are as follows – 'RecipeID', 'Name', 'Ingredients', 'Servings Per Recipe', 'calories', 'protein', 'carbohydrates', 'fat', 'saturated_fat', 'cholesterol', 'fiber', 'Sodium', 'Rating', 'Number of Ratings', 'URL'
- **Clustering:** To do clustering in these three datasets, we first calculate the quantities of nutrients in each recipe for a 1000 calories (this is done in order to make the recipes comparable). Once this is done, the recipes are then clustered into 3, 4 and 4 clusters respectively for breakfast, lunch and dinner using *K Means*. The elbow graph shown is for clustering in breakfast dataframe.



Dataset Preprocessing

- **User Dataset:** User profiles are being synthetically generated.
- **Nutrient Calculation:** The optimal nutrients for each user are calculated and appended to the row corresponding to the user in the dataset. This nutrient calculation has been done using BMR (basal metabolic rate) and physical activity level, and other formulas found upon research online.

| | personID | gender | age | height | weight | physical_activity_level | muscle | weightloss | hearthealthy | bmr | calories | protein | fat | fiber | carbohydrates | |
|--|----------|--------|-----|--------|--------|-------------------------|--------|------------|--------------|-----|----------|-------------|-------|-----------|---------------|------------|
| | 0 | 1 | 0 | 68 | 166 | 53.0 | 1 | 1 | 0 | 1 | 1157.512 | 1591.579 | 53.0 | 35.368422 | 22.282106 | 179.052638 |
| | 1 | 2 | 1 | 57 | 158 | 58.0 | 0 | 0 | 1 | 1 | 1300.041 | 1560.0492 | 46.4 | 34.66776 | 21.840689 | 175.505535 |
| | 2 | 3 | 0 | 24 | 184 | 65.0 | 0 | 1 | 1 | 0 | 1514.76 | 1817.712 | 52.0 | 40.3936 | 25.447968 | 204.4926 |
| | 3 | 4 | 0 | 49 | 178 | 69.0 | 2 | 1 | 1 | 0 | 1424.91 | 2087.49315 | 82.8 | 46.388737 | 29.224904 | 234.842979 |
| | 4 | 5 | 0 | 65 | 188 | 66.0 | 0 | 0 | 1 | 0 | 1358.869 | 1630.6428 | 52.8 | 36.236507 | 22.828999 | 183.447315 |
| | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| | 9995 | 9996 | 1 | 33 | 189 | 87.0 | 2 | 0 | 1 | 1 | 1973.571 | 2891.281515 | 104.4 | 64.2507 | 40.477941 | 325.26917 |
| | 9996 | 9997 | 0 | 77 | 189 | 74.0 | 3 | 0 | 1 | 0 | 1383.983 | 2145.17365 | 103.6 | 47.670526 | 30.032431 | 241.332036 |
| | 9997 | 9998 | 1 | 71 | 179 | 61.0 | 5 | 0 | 1 | 0 | 1361.533 | 2586.9127 | 109.8 | 57.486949 | 36.216778 | 291.027679 |
| | 9998 | 9999 | 1 | 33 | 157 | 55.0 | 4 | 1 | 0 | 1 | 1391.299 | 2399.990775 | 88.0 | 53.333128 | 33.599871 | 269.998962 |
| | 9999 | 10000 | 0 | 52 | 151 | 52.0 | 5 | 0 | 0 | 1 | 1171.075 | 2225.0425 | 93.6 | 49.445389 | 31.150595 | 250.317281 |



Reward Generation

- An arm is defined by a combination of clusters from breakfast, lunch and dinner (i.e we have $3 \times 4 \times 3 = 36$ combos or arms).
- For each user, we iterate over all arms (combos). Each arm is defined by a combination of clusters as mentioned above. So, arm 0 is defined by the cluster combo (0, 0, 0). Similarly, arm 1 is defined as (0, 0, 1), arm 2 as (0, 0, 2), arm 3 as (0, 1, 0) etc. We sample a recipe randomly from clusters i, j and k for an arm defined by (i, j, k) .
- We then calculate the reward as shown in the code on left.



```
x = cosine_sim(u, sample)
if (muscle == 1):
    diff = (sample[0]-u[0])/u[0]
    x = x + 0.5*diff
if (weightloss == 1):
    diff = (sample[3]-u[3])/u[3]
    x = x - 0.5*diff
if (heart == 1):
    diff = (sample[1]-u[1])/u[1]
    x = x - 0.5*diff
if (x > 1):
    x = 1
if (x<0):
    x=0
user_rewards_dict[(i, j, k)] = x
```


Reward Generation

| | (0, 0, 0) | (0, 0, 1) | (0, 0, 2) | (0, 1, 0) | (0, 1, 1) | (0, 1, 2) | (0, 2, 0) | (0, 2, 1) | (0, 2, 2) | (0, 3, 0) | ... |
|-------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----|
| 0 | 0.835199 | 0.557840 | 0.358004 | 0.046643 | 0.296029 | 0.000000 | 0.960872 | 0.551900 | 0.451923 | 0.472631 | ... |
| 1 | 0.545029 | 0.609466 | 0.348340 | 0.000000 | 0.054926 | 0.000000 | 0.434613 | 0.394769 | 0.206405 | 0.042048 | ... |
| 2 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | ... |
| 3 | 1.000000 | 1.000000 | 0.963419 | 1.000000 | 1.000000 | 0.983713 | 1.000000 | 0.934148 | 1.000000 | 1.000000 | ... |
| 4 | 0.974306 | 0.985010 | 0.985582 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 0.985262 | 1.000000 | 1.000000 | ... |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 0.538650 | 0.438592 | 0.236177 | 0.367275 | 0.371677 | 0.242172 | 0.493035 | 0.517051 | 0.322454 | 0.437758 | ... |
| 9996 | 0.875749 | 0.603778 | 0.802933 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 0.987667 | 1.000000 | 1.000000 | ... |
| 9997 | 1.000000 | 0.928468 | 1.000000 | 1.000000 | 1.000000 | 0.990706 | 1.000000 | 0.848566 | 1.000000 | 0.943023 | ... |
| 9998 | 1.000000 | 0.995334 | 0.595955 | 0.618899 | 0.773503 | 0.301467 | 1.000000 | 0.508460 | 0.287489 | 1.000000 | ... |
| 9999 | 0.980314 | 1.000000 | 0.831742 | 0.127738 | 0.432294 | 0.000000 | 0.861696 | 0.744734 | 0.747378 | 0.461207 | ... |

10000 rows × 36 columns



Context Formation

- In contextual bandits, context represents a combination of a user and an arm.
- User's context is defined by BMR, activity level, muscle goal, weight loss goal, heart healthy goal.
- Arms context is defined by the vector denoting nutrients for a particular combo.
- We concatenate the two vectors to form the context.

```
def make_context(user_df, cluster_combo_means, i, key):  
    user_features = [user_df.loc[i, 'bmr'], user_df.loc[i, 'muscle'],  
                    item_features = cluster_combo_means[key]  
    user_features.extend(item_features)  
    context = np.array(user_features)  
    if (np.dot(context.T, context)) > 1:  
        context = context/np.dot(context.T, context)  
    return context
```



LinUCB

- We first tried using Linear UCB, whose code is shown below.

```
def LinUCB(X, y, w, alpha):
    set1 = set([])
    for i in range(10, 1000):
        p = {}
        for key, value in cluster_combo_means.items():
            context = make_context(user_df, cluster_combo_means, i, key)
            p_t_a = np.dot(w.T, context) + alpha*np.sqrt(np.dot(context.T, np.dot(np.linalg.inv(np.dot(X.T, X)), context)))
            p[key] = p_t_a
        chosen_arm = max(p, key=p.get)
        set1.add(chosen_arm)
        print("chosen_arm = ", chosen_arm)
        r = rewards_df.loc[i, str(chosen_arm)]
        context = make_context(user_df, cluster_combo_means, i, chosen_arm)
        X = np.append(X, [context], axis=0)
        y = np.append(y, r)
        model = LinearRegression()
        model.fit(X, y)
        w = model.coef_
    return w, set1
```



Epsilon-Greedy

- We then used Epsilon-Greedy, whose code is shown below. We have taken epsilon to be 0.5.
- We chose Epsilon-Greedy over Linear UCB because we realised that even upon experimenting with various values of alpha, the algorithm was not exploring all the arms which is undesirable.



```
def epsilon_greedy(X, y, w, epsilon):
    set1 = set([])
    for i in range (0, 10000):
        random_number = np.random.random()
        if (random_number < epsilon):
            chosen_arm = random.randint(0,35)
        else:
            p = {}
            for key, value in cluster_combo_means.items():
                context = make_context(user_df, cluster_combo_means, i, key)
                p_t_a = np.dot(w.T, context)
                p[key] = p_t_a
            chosen_arm = max(p, key=p.get)
        set1.add(chosen_arm)
        r = rewards_df.loc[i, str(chosen_arm)]
        context = make_context(user_df, cluster_combo_means, i, chosen_arm)
        context = np.squeeze(context)
        X = np.append(X, [context], axis=0)
        y = np.append(y, r)
        model = LinearRegression()
        model.fit(X, y)
        w = model.coef_
    return w, set1
```


Results

- By the end of the run of the algorithm, if for a user, we are doing exploitation, the best arm should be picked upon exploitation.
- This means that, that particular arm should have the highest reward in the rewards dataframe for that user.
- It can be seen that at the 997th step, the arm chosen upon exploitation (0), is 12. And in rewards dataframe, the rewards for that arm is 1.

```
rewards_df.loc[997, '12']
```

1.0

```
w, set1 = epsilon_greedy(X, y, w_initial, 0.5)
```

```
0 12
1 33
0 12
0 12
0 12
0 12
0 12
0 12
1 4
1 0
1 11
1 5
1 21
1 0
1 18
0 12
1 31
0 12
1 17
1 12
```




Challenges Faced

- The main challenge faced was how to form the context vector to make it learnable. We tried various methods for this. Element wise addition of user and arm vector, tensor product between user and arm vector, Hadamard Product between the user and arm vector. But all these methods did not improve the performance of the model, it was almost same as the performance as concatenation of the two vectors.
- Defining rewards was a difficult task. We first tried the converse of the norm of the difference between user and sample vector. Then we went on to do the similarity metric with some noise so as to make the rewards stochastic. Finally, random sampling made the most sense since in real life too people would pick random recipes.
- Clustering recipes on what basis was also a challenge.
- We also faced many errors initially during algorithm implementation. Most of them were input shape related.



Approach - 2

Recommendations, Ratings, Epsilon
Greedy, Decaying epsilon



Data Preprocessing





Filtering According to Diet Goal

Filtering recipes based on diet goals by setting appropriate thresholds on nutrient values.



Muscle Gain

Focus on high protein intake.



Heart Healthy

Promote foods low in saturated fats and rich in fruits, vegetables

Muscle Gain

- Below are the constraints for filtering recipes from the dataset for users who want to gain muscle.

```
muscle_breakfast = breakfast_df[  
    (breakfast_df['Protein'] >= 20) # High protein content  
]
```

```
muscle_lunch_df = lunch_df[lunch_df['Protein'] >= 20]
```

```
muscle_dinner_df = dinner_df[dinner_df['Protein'] >= 30]
```

[Reference for above thresholds](#)



Heart Healthy

- Below are the constraints for filtering recipes from the dataset for users who want to prioritize heart health.

```
heart_breakfast_df = breakfast_df[(breakfast_df['Total Fat'] <= 8) &  
                                   (breakfast_df['Cholesterol'] <= 20) ]
```

```
heart_lunch_df = lunch_df[(lunch_df['Total Fat'] <= 15) & (lunch_df['Cholesterol'] <= 20) ]
```


```
heart_dinner_df = dinner_df[(dinner_df['Total Fat'] <= 15) &  
                             (dinner_df['Cholesterol'] <= 20) ]
```

[Reference for above thresholds](#)



Clustering

- We did clustering for each of the filtered data frames using **k-means++**.
- And then choose the optimal number of clusters using elbow method.
- According to the graph choosing **3** as the optimal clusters and applying k-means++ with number of clusters as 3



Cluster Centers:

```
[[ 35.48281093 154.99727481 31.22710167 27.14692591]
 [ 59.43542627 67.42255599 56.24081259 12.91914758]
 [ 23.62900045 230.33610679 5.40671645 43.75584999]]
```





Greedy Epsilon MAB

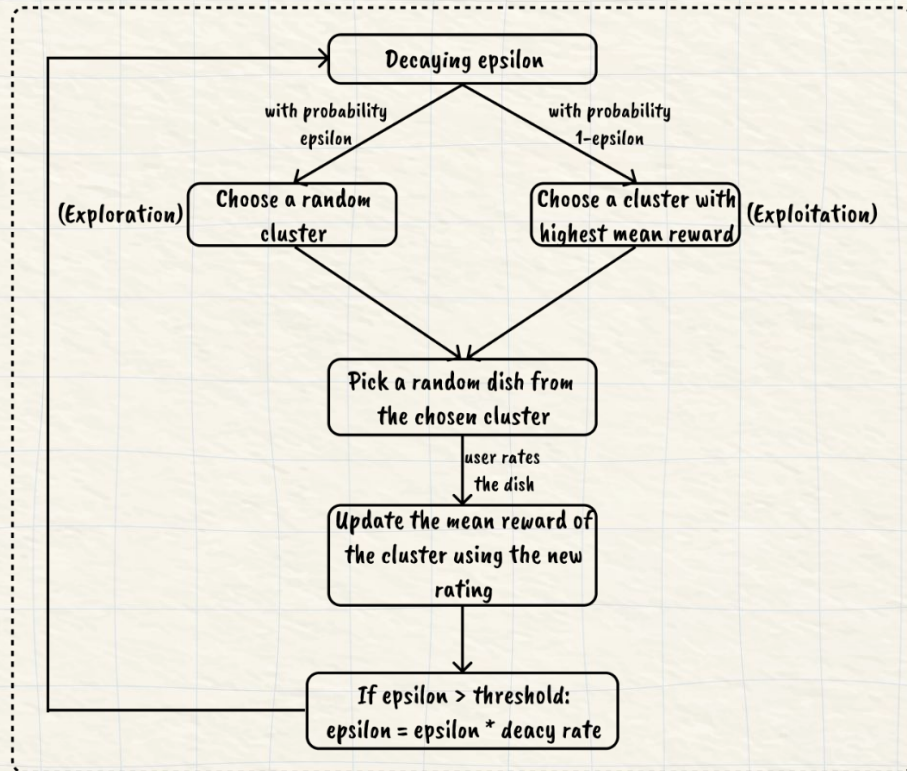


New User

Enters Goal

Goal specific
dataframe chosen

Greedy-epsilon MAB

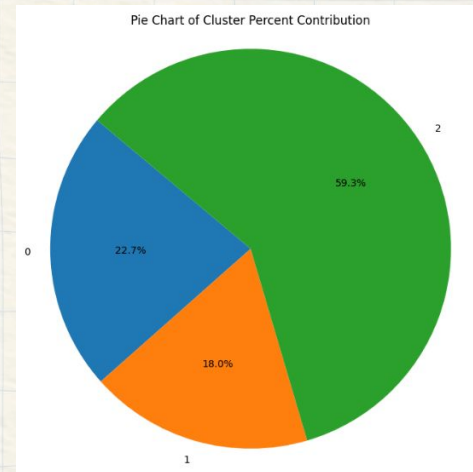
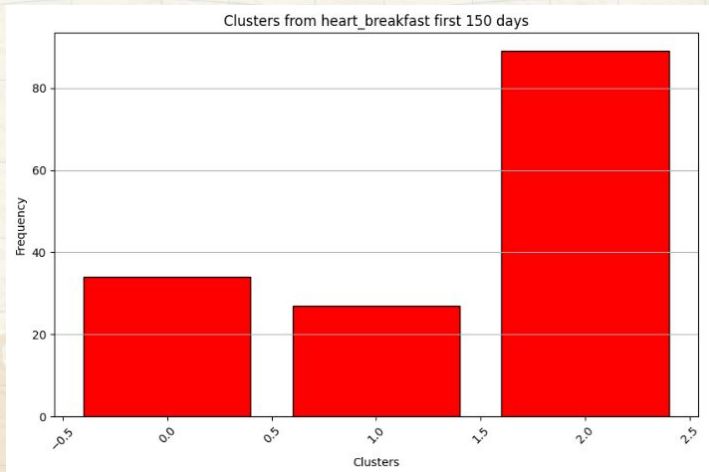


Results

To test run our system, we have generated 365 random ratings that would be given for breakfast recommendations of each day for 365 days. Below are the visualizations depicting the frequency of chosen arms (clusters).

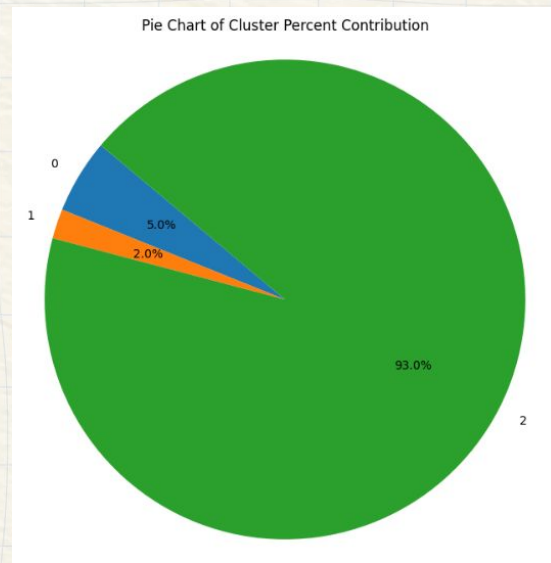


1. For first 150 days.




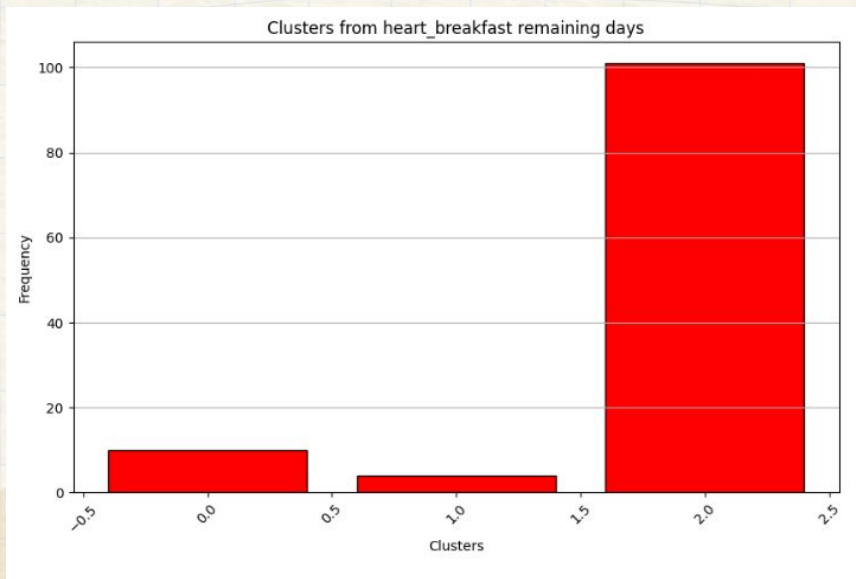
Results

2. For Next 100 days.



Results

3. For Remaining days.



Challenges & Improvements



1. Clustering of recipes: Clustering the recipes was a challenge as we wanted the clustering to be relevant to the taste of the food item. According to the research paper ^[1] nutrients like protein, carbohydrates, fats, etc influence the taste of the food item. Hence we applied kmeans on the nutrient values after scaling them to per 1000 calories.

2. Non desirable cluster being chosen majority of times during exploration initially: It so may happen during the initial phase where exploration is given priority, a specific cluster might get recommended everytime randomly and even though the user doesn't prefer that cluster, the same cluster would be recommended during exploitation too as others mean would still be 0. Hence to tackle this initial mean ratings have been set to 2 instead of 0.

3. Epsilon becoming too small: Due to decaying epsilon, epsilon may become too small and exploration would never happen. This would lead to similar recommendations from there on. Hence we have put a lower bound on the value of epsilon such that it won't decay below the lower bound so that a minimal amount of exploration would always persist.

[1] van Langeveld, Astrid & Gibbons, Shannon & Koelliker, Yvonne & Civile, Gail & de Vries, Jeanne & Graaf, Cees & Mars, Monica. (2016). The relationship between taste and nutrient content in commercially available foods from the United States. Food Quality and Preference. 57. 10.1016/j.foodqual.2016.10.012.



Deployment on Streamlit



Automatic Diet Recommendation x +

localhost:8501/ContextualDietRecommendation#diet-recommendation-system

Deploy

Diet Recommendation System

Welcome

ContextualDietRecommendation

Age

21 - +

Height(cm)

174 - +

Weight(kg)

62 - +

Gender

☒ Male

☐ Female

Choose your activity level:

Exercise 4-5 times/week

Choose your weight loss plan:

Muscle Gain


Generate

BMI CALCULATOR

Body Mass Index (BMI)

20.48 kg/m²

Normal



Thanks!

